



# From the other side of the Atlantic, How to improve upon the MU5 design.

James C. Brakefield

## Abstract:

The MU5 computer and its commercial counter-part the ICL2900 have become historically interesting in much the same way that Algol 60 prompted research in programming languages. What follows is my extrapolation of where the MU5 architecture leads.

## Preliminaries:

The idea undergoing evolution is that of what a memory address should be. Ie., the earliest developments where those of indirect addressing and index registers. Since then there has been continuous experimentation with greater and greater variety of addressing modes. On those computers whose word size is greater than the address size there has been experimentation with additional fields associated with the address qualifying the object referenced by the address. These kinds of address words are called descriptors.

A somewhat parallel evolution is to add non-data bits to the word to qualify the word by itself. These are called tags. There is a practical difficulty with tags in that they are inherently less memory efficient than descriptors and they do not allow qualification of less than a whole word. Finally, some computers provide both tags and descriptors.

The reason tags are inefficient is that every word must have a tag. If the word contains fields, the tag can not fully describe them. This limits the use of fully packed records. If several words all have the same tag, such as in an array, then their tags are all the same. This again is inefficient.

What makes this all the more interesting is the use of virtual memory. With large virtual address spaces complete files can be mapped into the address space of an application program. The record structure of the file can then be exploited in the program. Computer hardware which facilitates record manipulation in this environment is very handy.

This paper takes what I call an American viewpoint and restricts descriptors to describing only single objects. By this I mean that vectors and arrays are not supported directly with corresponding descriptors. The problem is that current computer architecture uses addresses of about the same size as words. Thus, to include a length or separation field in the descriptor requires a double-precision (ie. two word) descriptor.

The most efficient approach then seems to be that single as opposed to vector operand instructions use single precision descriptors and if the computer has vector instructions, this is where to use double precision descriptors.

The architecture:

The operands for instructions shall be unsigned integers, signed integers, floating point numbers, and collections of these. The memory allotment for such operands shall be variable. Unsigned integers shall be allowed to have any power-of-two size from one to 64 bits. Signed integers shall be allowed power-of-two sizes from two to 64 bits. Floating point numbers shall have sizes from four to 64 bits. A given implementation may provide additional sizes at the expense of a smaller address space.

The choice of formats for integers and floating point is not entirely open. Conversion between signed and unsigned is simplest if two's complement is used. Extraction of signed integers from memory and the extension of short integers to longer integers favors sign and magnitude with sign at the least significant end. The current popularity of two's complement will probably ensure its survival.

With floating point the extraction of the smaller formats from memory and extension of short immediate formats argue for having the number of exponent bits proportional to size. One easy way to do this is to interleave them with the mantissa bits on say a one-to-three ratio. This requires a reformatting circuit between the registers and the floating point arithmetic unit, which may be the same unit doing integer arithmetic.

Collections of data are records. Record sizes have no upper limit. Thus as a minimum records of 4, 8, 16, 32, 64, ... bits must be supported.

Instructions shall be coded in 8 bit multiples. Descriptors referencing instructions will not be required to specify size.

The operand type and size is always coded with the address. This becomes the descriptor. A complete list of data structures is:

- The unsigned integer of 1, 2, 4, 8, 16, 32, & 64 bits
- The signed integer of 2, 4, 8, 16, 32, & 64 bits
- The floating point number of 4, 8, 16, 32, & 64 bits
- The record of 4, 8, 16, 32, 64, ... bits
- The descriptor of 32 and possibly 64 bits
- Program code in multiples of 8 bits

- The first option is to support 3x sizes, ie. 3, 6, 12, ...
- The second option is to support 5x sizes, ie. 5, 10, 20, ...
- The "n"th option is to support all sizes from 1 to 64 bits

On implementations which use a 3x word size (such as 24 and 48 bit word size), both the 1x and 3x sizes must be supported. In this case, the required descriptor size will be a 3x size. Other descriptor sizes may be optionally supported.

The 1x descriptor:

There are two fields with leading one-bit coding. The type field (least significant part of the descriptor) is coded as follows:

xxxxxxx1	unsigned integer
xxxxxxx10	signed integer
xxxxx100	floating
xxxxx000	record

And the address/size field:

yyyyyy11	unsigned integer of one bit
yyyyyy1xx	unsigned integer or signed integer of two bits
yyyyylxxx	unsigned, signed, float, or record of four bits

.....

Examples:

yyyyyl001	(unsigned of four bits)
yyyl0010	(signed of eight bits)
yy1000000	(record of 32-bits)

The 3x descriptor:

The least significant bit is a 3x specifier. If set, size is multiplied by three. There is one less address bit.

xxxxx10	unsigned integer of 1, 2, ... bits
xxxxx11	unsigned integer of 3, 6, ... bits
	signed, floating, & records similarly

The 3x-5x descriptor:

The 3x descriptor and 1x descriptor only allow an item of 2\*\*n size to be on a 2\*\*n boundary. This is a little restrictive in the larger sizes. The 3x version takes one bit, the 3x-5x version takes two bits, but there is enough codes so that an item of 2\*\*n size can be on a 2\*\*n-1 boundary.

xxxxxxx0	of 2, 4, ... bits on 1, 2, ... boundary
xxxxxx010	unsigned
xxxxxx100	signed
xxxxx1000	float
xxxxx0000	record
xxxxx0110	bit on bit boundary
xxxxx1110	unused
xxxxxxx01	of 3, 6, ... bits
xxxxxx101	unsigned
xxxxx1001	signed
xxxx10001	float

xxxx00001	record
xxxxxxx11	of 5, 10, ... bits
xxxxxx111	unsigned
xxxxx1011	signed
xxxx10011	float
xxxx00011	record

The instruction set:

The use of instructions with fields for operand specification is abandoned. The instruction string consists of intermixed operand specification and instruction commands. Operand specification may be arbitrarily complex, however, such complexities do not interfere with virtual addressing requirements.

The result is that instruction commands are function calls with little variability. The complexity is in operand specification which is completely removed from command specification.

Another consideration is to the register and stack facilities of a given implementation. The author's preference is to provide a limited number of general purpose registers (eight will do) and to emphasize the use of a stacking mechanism. There are to be two stacks. One holds intermediate calculations, descriptors, and subroutine parameters. The other holds return addresses, loop counters and limits, and any other control information (status bits, etc).

With this much specified, what remains is determination of the addressing modes provided. There are several general categories:

- A) Operand comes from the instruction stream (immediate data)
- B) Operand comes from one of the registers or one of the stacks
- C) Operand is reached indirectly through a descriptor in one of the registers or from within one of the stacks
- D) The descriptor in C is modified by indexing, offsetting, or indirection
- E) The offset or index is arrived at by the general operand specification process
- F) The descriptor is pre-decremented or post-incremented by the size of the operand

Due to the variable length nature of operand specification it preceeds the instruction which utilizes the operand. The situation is essentially that of "reverse polish notation". Any operands needed by an instruction, but not specified are taken from the value stack.

A mode:

Immediate operand of 0..63

Immediate operand of 8n bits where n = 1..8

B mode:

Register operand, register 0..7  
Top of value stack  
Top of control stack (typically for loop count)  
Operand is at an offset from one of the address registers:  
either stack pointer, base or global pointer, frame  
pointer, offset  
value specification follows in operand format

C mode:

Register holds descriptor of operand

Top of value stack holds descriptor of operand

A or B mode operand is descriptor of operand instead of  
operand itself

D mode:

C mode descriptor is indexed by signed integer (index is  
multiplied by the size of the operand and added to the address)

C mode descriptor is offset by a second descriptor, the size  
of the operand is taken from the second descriptor, the offset  
must be less than the size of the original descriptor operand

The C mode descriptor is taken as an indirect address of yet  
another descriptor

E mode:

The D mode index and offsets are specified by following  
operand specifications

F mode:

Mode flag preceeds operand specification and identifies that  
pre-decrement/post-increment will be applied

All of the above modes can be combined to use less than 256  
codes. The remaining codes are used for instructions (op codes).  
An instruction such as ADD comes in three forms: single operand  
(increment), two operand (replace add), and three operand (add  
and store). Unspecified operands come from the value stack. ADD  
would also have an integer/floating point mode in each of the  
three forms. Register and stack references assume the operand is  
in the correct form (signed, unsigned, floating, or descriptor).  
References through descriptors use the form specified by the  
descriptor which may require automatic type conversions to match  
the form specified by the op-code.

Moves & format changes

Move, Move unsigned, Move signed, Move floating, Move descriptor  
Exchange, Clear

Unsigned, Signed, & Floating

Test, Compare, Add, Subtract, Multiply, Divide, Abs, Minus, Max,  
Min

Conditional branches, Looping

Branch, Jump, Jump & save, Count up/down loop, Count to limit  
loop, Case

Logical

And, Or, Exclusive or, Not, Masked test

Stack operations  
Pop, Duplicate, Rotate up/down

References:

1. Morris, D. and Ibbett, R.N., 1979, The MU5 Computer System, Springer-Verlag, N.Y.
2. Bishop, J.M. and Barron, D.W., Aug. 1981, Principles of Descriptors, The Computer Journal, 24:3:210-221.
3. Izatt, W.T. and Schmitz, E.A., Nov. 1981, Data structures and descriptors in the ICL2900 series and beyond, The Computer Journal, 24:4:301-307.
4. Ibbett, R.N. and Capon, P.C., Jan. 1978, The Developement of the MU5 Computer System, CACM 24:1:13-24.