

Use of Neural Models for Cognitive Processing

2004 Sandia-UNM Cognitive Systems Workshop

James C. Brakefield © 2004

Abstract:

Results of a feasibility study show that system level neuron and synapse models can be implemented in affordable scalable real-time systems. Current limitations in scientific knowledge make it impossible to accurately model complete mammalian nervous systems. However, cognitive systems can be designed and tested without biological details, and therefore allow an avenue of research using neurological based cognitive processing. Example pulse mode neurological models and implementation architectures will be presented.

Contents:

Biases and Prejudices
Diffuse Neuron Projections
Gogli Image
Brain Simulation Issues
Neuron Firing Patterns
System Simulation “Rational”
1000X Reduced Configurations
Human Brain 101
Mammal Brain Sizes
Implementation Choices
Sample Compute Module

Micro-Electronics Background
Engineering Considerations
Hardware Simulation Alternatives
Neuronal Modeling
Neural Biology Web sites
Biological Connectivity Problem
Engineering Risk Assessment
Digital/Hardware Learning
Additional References
Gogli Images 2 & 3

Use of Neural Models for Cognitive Processing (cont'd)

Biases and Prejudices - 1

Quantum effects don't matter:

Anthony J. Bell: *Levels and loops: the future of artificial intelligence and neuroscience*, Phil. Trans. R. Soc. Lond. B (1999) 354, 2013-2020.

Consciousness irrelevant:

Look to Gazziniga's interpreter to understand consciousness:

M.S. Gazzaniga: *The Social Brain: Discovering the Networks of the Mind*, 1985, Basic Books. <http://www.dartmouth.edu/~cogneuro/Gazzaniga.html>

Intelligence a function of scale and architecture and chemistry

Scale: Typical neurons have 6000 inputs & outputs

Human brain: 10^{10} - 10^{12} neurons

Architecture: More than just a uniform neo-cortex surface

Chemistry: Diffuse projections: dopamine, serotonin, noradrenaline, acetylcholine

Use of Neural Models for Cognitive Processing (cont'd)

Biases and Prejudices - 2

Scale matters

Lesson from Mandelbrot: new tools lead to new science:

Computer graphic workstation allowed visual investigation of fractals:

Benoit B. Mandelbrot: The fractal geometry of nature, 1977, W.H. Freeman.

Lack perspective for devices with thousands of inputs and outputs.

Sparse Distributed Memory (SDM) and others typically use 500+ inputs.

<http://www.rni.org/kanerva/homepg.html>

Usual methodology for high In/Out counts: Operate in time domain

Usual methodology for high unit numbers: Analyze in continuous domain

Izhikevich E.M. (2005) *Dynamical Systems in Neuroscience: The Geometry of Excitability and Bursting*.

Use of Neural Models for Cognitive Processing (cont'd)

Biases and Prejudices - 3

Timing matters

A general purpose simulator must handle axon and dendrite transmission delays

Izhikevich E.M., Gally J.A., and Edelman, G.M.: *Spike-Timing Dynamics of Neuronal Groups* (2004) Cerebral Cortex, in press.

Spiking models now proven necessary to match biological data:

Simon Thorpe: *Ultra-rapid scene categorization with a wave of spikes*.
(150 ms to do image processing, or ~10 ms per neural stage)

Jacques Sougne & Robert French: *Synfire Chains and Catastrophic Interference*,
Proceedings of 23rd Annual Conference of the Cognitive Science Society.
(millisecond accuracy of neuron firing after a delay of 300ms)


Noise matters

Behavior can depend upon a single spike (also Thorpe)

Real-Time Digital Simulation of the Human Nervous System (cont.)

Diffuse Neuron Projections

Kenji Doya: <http://www.his.atr.jp/~doya/topics.html>



Use of Neural Models for Cognitive Processing (cont'd)

Biases and Prejudices - 4

System level modeling

Must do science on the whole to completely understand the whole:

Some ~50 brain regions with extensive interconnections, so each region in isolation presents an incomplete picture.

Cortex exhibits same phenomena in even greater complexity.

The brain is an engineered machine, not a general-purpose computer.

Attempts to use neural models to engineer cognition somewhat like early history of airplanes.

Meta Learning based on Diffuse Projection Neurons


ATR Computational Neuroscience: <http://www.cns.atr.jp/top.html>

Kenji Doya: <http://www.his.atr.jp/~doya/topics.html> !!!

Doya K. (2002). Metalearning and neuromodulation. Neural Networks, 15, 495-506.

Real-Time Digital Simulation of the Human Nervous System (cont.)

Golgi Atlas of the Postnatal Mouse Brain: www.cajal.csic.es/valverde



Use of Neural Models for Cognitive Processing (cont'd)

Brain Simulation Issues

General purpose models:

Some three to four thousand distinct neuron types of operation and morphology in mouse brain, most of which have not been studied as yet.

A parameterized model which encompasses most neural types of operation:

Eugene M. Izhikevich: *Which Model to Use for Cortical Spiking Networks*, IEEE. Trans. On Neural Networks (2004, submitted). Also: *Simple Model of Spiking Network*, ibid. (has Matlab & C++ source code).

Connections have higher information content than synapse state.

Log base 2 of number of potential synapse inputs is ~34 bits

Log base 2 of number of synapse configurations < 8 bits

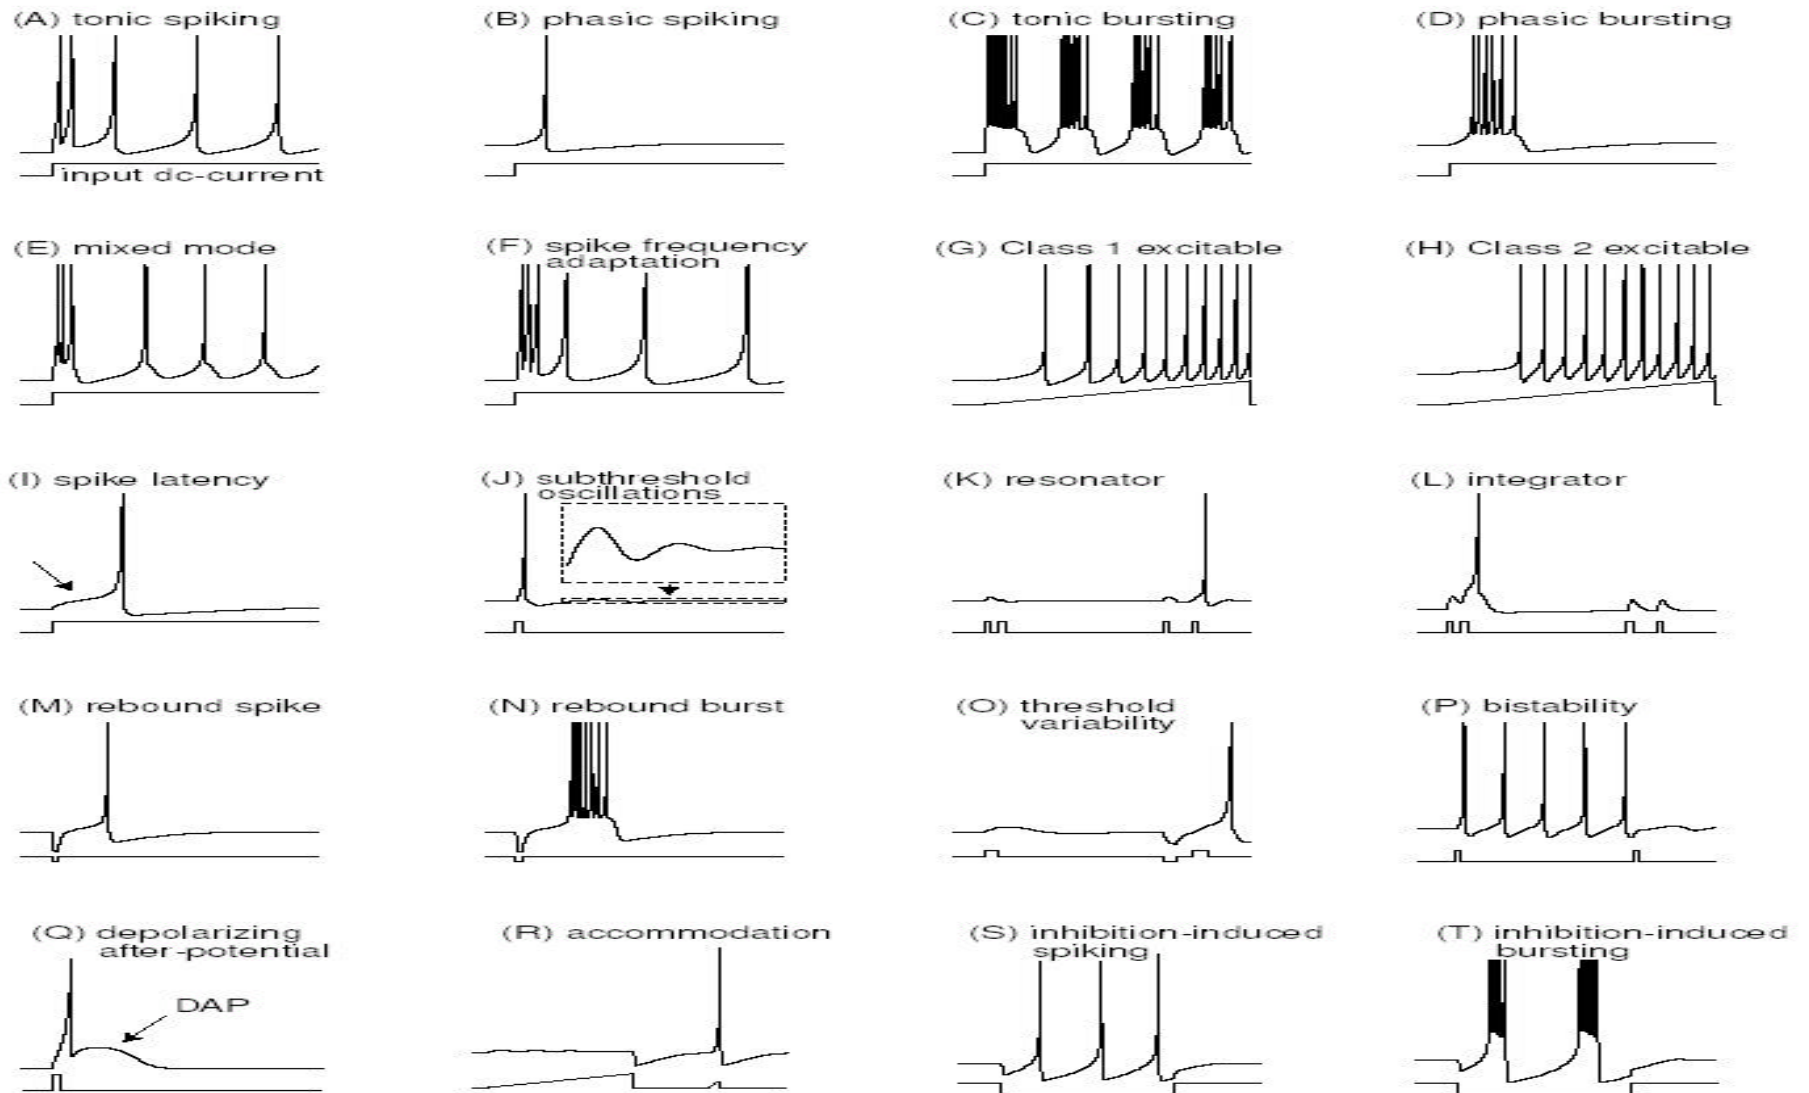
Log base 2 of synaptic delay ~6 bits (40ms at 0.5ms resolution)

$33 \gg 8 + 6$

Also holds true for digital circuits: ASICs (gates) or FPGAs (lookup tables)

Use of Neural Models for Cognitive Processing (cont'd)

Neuron Firing Patterns E.M. Izhikevich: *Which Model to Use for Cortical Spiking Neurons*



Use of Neural Models for Cognitive Processing (cont'd)

System Simulation “Rational”


Real-Time Digital Simulation of the Human Nervous System

The objective is to study the feasibility of a real-time system level simulation of the human brain. The result is that it is considered feasible using FPGA and DRAM technology at a chip cost of \$100+ Million US dollars.

Why even consider such an ambitious project?

- A) The study itself adds insight into current knowledge of the brain.
- B) It provides a framework from which to organize existing knowledge.
- C) It forces component level models of neurobiology to be considered from the system or whole.
- D) Silicon costs are expected to continue to decline. In a decade, the issue will not be why, but when and by whom. (see www.ad.com/tech.html, www.cybernetics.demon.co.uk/Plan.html, www.artificialbrains.com)
- E) A configuration reduced by one thousand times, is suitable for the system level simulation of small mammals. Silicon cost also reduces by one thousand, leading to a tool affordable by many research institutions.
- F) The reduced configuration is compact and therefore may have military use.**
- G) The system level perspective puts new light on what is overlooked by researchers not interested in system level issues.**
- H) A facility and the reduced scale facility are useful for many other kinds of complex system modeling.**

Real-Time Digital Simulation of the Human Nervous System (cont.) Reduced Configurations (~1000X reduction in size)



Simplified Human



Small Mammal



Military Applications

Use of Neural Models for Cognitive Processing (cont'd)

Human Brain 101

Basic Facts (<http://faculty.washington.edu/chudler/facts.html>):

1,300-1,400 grams:	$1.3 \times 10^6 \text{ mm}^3$	
60% gray matter:	$0.8 \times 10^6 \text{ mm}^3$	(gray matter is neurons, white matter is wiring)
cortex surface area:	$0.20\text{-}0.25 \times 10^6 \text{ mm}^2$	(cortex is largest part of human brain)
cortex thickness:	average is 2.0 mm,	range of 1.5 to 4.5 mm
neurons in cortex:	$10\text{-}20 \times 10^9$,	100K-150K per mm^2
synapses in cortex:	$60\text{-}240 \times 10^{12}$,	$10^8\text{-}10^9$ per mm^3
Koch:	page 87: average of 6×10^8 per mm^3 ,	up to 10^9 per mm^3
Shepherd:	page 471 (Beaulieu & Colonnier):	2.78×10^8 per mm^3
Sejnowski:	page 51:	10^9 per mm^3 , recently 2×10^9 per mm^3
(Cortex dominates brain, cerebellum has $\sim 30 \times 10^9$ neurons and only $\sim 6 \times 10^{12}$ synapses)		
neuron firing rate:	1-200 Hertz,	resting: 1-10 Hertz, visual cortex: ~ 40 Hertz


Shepherd: The Synaptic Organization of the Brain, 4th ed., 1998, Oxford University Press

Koch: Biophysics of Computation: Information Processing in Single Neurons, 1999, Oxford Press

Churchland & Sejnowski: The Computational Brain, 1992, MIT Press

Use of Neural Models for Cognitive Processing (cont'd)

Mammal Brain Sizes, From Zhang & Sejnowski, PNAS 97:10:5622



Use of Neural Models for Cognitive Processing (cont'd)

Implementation choices


Using 3D torus interconnect of identical processing units

- 1) Microprocessor farm (e.g. Red Storm).
- 2) Micro-controller + DRAM (~1GB for synapse data) + FPGA (for store & forward network). (e.g. ARM microprocessor & memory on small circuit board)
- 3) FPGA (with on chip micro-processor) + DRAM (~1GB).
- 4) Synchronous simulation: FPGA + DRAM (~1GB).
- 5) Asynchronous simulation: FPGA + DRAM (~1GB) + dual-port SRAM (event queues).

In all cases small FPGA or portion of FPGA or an ASIC needed for 3D torus interconnect.

Use of Neural Models for Cognitive Processing (cont'd)

Sample Compute Module: www.simtec.co.uk/products/EB110MOD/



Use of Neural Models for Cognitive Processing (cont'd)

Micro-Electronics Background

AMD “Hammer” Opteron PC Chip: 2-3Ghz, 64-bit address space,
“Red Storm”: http://www.cray.com/news/0210/sandia_redstorm.html, 10K chips,
3D mesh inter-connect, for Sandia National Laboratories

Samsung “Halla” ARM Chip: ~600Mhz, 32-bit address space, integer only, also Intel PXA27x
http://www.convergencepromotions.com/pdf/Samsung_article.pdf

Intrinsity *Fast*MATH signal processor: 2Ghz 32-bit processors, 17 per chip
http://www.intrinsity.com/products/products_fastmath_set.htm

FPGAs with 200-300 Mhz clock rates: Xilinx (www.xilinx.com)
Spartan-3 XCS2000: 565 I/Os, **\$25** each in high volume, Has:
DDR (double data rate) I/Os, (40) 18K dual port RAMs, (40) 18x18 multipliers

Other FPGAs: Altera(www.altera.com) & Lattice Semiconductor(www.latticesemi.com)
Altera: Stratix & Stratix II: similar to Spartan-3 and/or Virtex II
Xilinx: Virtex II Pro has embedded PowerPC processors

RLDRAM: Commodity DDR DRAM: ~\$0.10 per megabyte
Dual-Port RAM: Cypress and IDT: up to 18 Mbits, up to 72-bit I/O.

Use of Neural Models for Cognitive Processing (cont'd)

Engineering Considerations

Synapse state: 48 bits per synapse:

Connectivity: if 10^{10} neurons, then 33 bits per synapse for source identification

Because axon arbors are compact,

Half of connectivity handled by routing, and

Other half by RAM address (~16 bits)

Aggregate delay: 8 bits

Strength and type of synapse: 8 bits

State: 8 bits, double buffered (“current” and “update” buffers)

FPGA chip + DRAM

~ 2 x 2 inch circuit board

FPGA on one side, ~16 DRAM chips on other (+ dual-port RAM for async simulation)

400-600 Mhz DRAM data rate on 480 FPGA pins

Processing at 200-300 Mhz rate on each of 20 parallel streams (for sync simulation)

~25,000 neurons per circuit board at a 40 Hertz update rate

May need a heat sink (~10 watts per circuit board)

Interconnections

Flattened 3D torus with 200 Mhz serial links between nodes

Part of each FPGA chip devoted to store & forward mechanism

For 10^{14} synapses (a low-end estimate) & 40 Hertz synchronous update rate:

10^6 modules:

10^6 FPGAs: \$25 M

6×10^{14} bytes DRAM: \$60 M

Use of Neural Models for Cognitive Processing (cont'd)

Hardware Simulation Alternatives - 1

Asynchronous simulation

~ 1ms time step

Computational effort is lower than with synchronous simulation

Requires dual-port RAM for event queues

- One port to post events, driven from incoming and outgoing spikes

- One port to process events

- Higher cost & circuit board area due to dual-port RAM chip(s)

DRAM Memory used mostly in random access mode, so high memory bandwidth not possible

Architectural parallelism:

- 1) Store & forward network

- 2) Incoming spike processor

- DRAM operated in pipeline mode (for synapse examination)

- 3) Event processors

- Four way parallel processing possible

- DRAM operated in random access mode

Use of Neural Models for Cognitive Processing (cont'd)

Hardware Simulation Alternatives - 2

Synchronous simulation

Minimum update rate of 40Hz, 100Hz affordable

Memory bandwidth requirement scales with update rate

To increase update rate

 Increase number of FPGAs

 Increase DRAM data rate

 Total DRAM stays the same

Memory accessed in pipeline (streaming) mode

High parallelism is pin limited

Conversion of asynchronous models to synchronous implementation:

 Spike activity in previous cycle encoded into ~8-bit representation

Architectural parallelism:

 1) Store & forward network

 2) Synapse processors (~20 operating at 200 Mhz)

 DRAM interface (480 bits wide operating at 600 Mhz?)

Each incoming axon has an on-chip RAM word for “spike” representation

Use of Neural Models for Cognitive Processing (cont'd)

Hardware Simulation Alternatives - 3

Experiment Process Flow

Given the connectivity and model parameterization:

- 1) Derive axon, synapse, spine & dendrite-tree delays
- 2) Neurons allocated to FPGA chip/modules topographically
- 3) Parallelism requires allocation of synapses and/or neurons to various pipes
- 4) Memory locations allocated
- 5) Axon destination ranges (voxels on the 3D torus) allocated
- 6) FPGA chips configured
- 7) DRAM initialized
- 8) Run made
- 9) DRAM contents saved for analysis

Use of small scale Processor Farm (e.g. Red Storm) for support purposes:

Disk drives to load/save DRAM

Monitor runs

One processor per ~1K FPGA modules?

Use of Neural Models for Cognitive Processing (cont'd)

Neuronal Modeling

System Level Neuron Models:

Izhikevich: <http://www.nsi.edu/users/izhikevich/publications/>

Mathematician, Matlab and C++, Single parameterized model

Maass: <http://www.igi.tugraz.at/maass/Welcome.html>, <http://www.lsm.tugraz.at>

Computer Scientist, Matlab and C++, Pulse mode, adaptive

Maass & Bishop eds: Pulsed Neural Networks, 1998, MIT Press

Koch: <http://klab.caltech.edu/index.shtml> Neurobiologist

Koch & Segev eds: Methods of Neuronal Modeling, 2nd ed, 1998, MIT Press

Hopfield: <http://neuron.princeton.edu/~moment/Organism/index.html>

Physicist, Matlab

Van Essen & Anderson: <http://stp.wustl.edu/>

Neurobiologist, NESim written in Matlab, Java user interface

Thorpe, Delorme & VanRullen: SpikeNET: <http://www.spikenet-technology.com>

French/USA reseachers, SpikeNET does fast image recognition

Doya: <http://www.his.atr.jp/~doya/topics.html>

Mathematician, Matlab, meta-learning, e.g., feedback mechanisms

Use of Neural Models for Cognitive Processing (cont'd)

Neural Biology Web sites

Boston University & Harvard & others: <http://synapses.mcg.edu/>

Reconstruction from electron microscope sections

3D animations(from reconstructions), electron-microscope pictures, really great web site !!

Instituto Cajal: <http://www.cajal.csic.es/valverde/>

Web site has three Gogli method slices of postnatal mouse, for more slices:

Facundo Valverde: Gogli Atlas of the Postnatal Mouse Brain, 1998, Springer-Verlag

See also: George Paxinos & Keith Franklin: The Mouse Brain in Stereotaxic Coordinates, Deluxe Second Edition (with CD-ROM), 2000, Academic Press.

Cajal Neuroscience Research Center: <http://marlin.life.utsa.edu/>

Has axon arbors as well as dendrite trees

Bennett & Wilson: “Synaptology and Physiology of Neostriatal Neurons” in: Brain Dynamics and the Striatal Complex, Miller & Wicken eds, in press, Harwood Academic

Dissection: <http://www.vh.org/adult/provider/anatomy/BrainAnatomy/BrainAnatomy.html>

Mammal brain collections: <http://brainmuseum.org/>

Mammalian cortex sizes: Zhang & Sejnowski, “A universal scaling law between gray matter and white matter of cerebral cortex”, PNAS, May 9, 2000, vol.97, no.10, pgs 5621-5626 (on the web).

Real-Time Digital Simulation of the Human Nervous System (cont.) Biological Connectivity Problem

Two automated approaches to tracing neurons:

Knife Edge Scanning Microscopy:

Wonryull Koh & Bruce McCormick: <http://research.cs.tamu.edu/bnl>

Specifications for Volume Data Acquisition in Three-Dimensional Light Microscopy
500nm x 200nm x 200nm voxels, 200 Mhz pixel rate.

Two photon laser scanning microscopy:

“All-Optical Histology Using Ultrashort Laser Pulses”, Neuron, vol. 39, pp. 27-41, 2003,
Tsai, Friedman, Ifarraguerri, Thompson, Lev-Ram, Schaffer, Xiong, Tsien, Squier & Kleinfeld.
<http://www.tsienlav.ucsd.edu/Publications>

Neither has 100nm resolution (dendrite spines are about 100nm in diameter).

Effort to determine neuron type from biochemistry:

Paul Allen’s Brain Atlas: <http://www.brainatlas.org/>

Use of Neural Models for Cognitive Processing (cont'd)

Hardware Simulation Alternatives - 4

FPGA implementation superior to microprocessor implementation

Greater range of implementation choices

- Data width adjustable

- Pipelined processing

High speed self clocking serial interfaces readily available (800 Mhz and higher)

Parallelism readily available

Greater memory bandwidth (500 signal pins typical)

Same chip does store & forward function as well as processing

For digital circuits the information content is in the connections

ASICs: 2 and three input gates, thousands to millions of possible inputs

- Three bits to identify gate type, 20 bits per input to identify each input

- E.g., 40 bits for inputs, 3 bits for logic

FPGAs: 3, 4, 5, 6 input lookup tables (LUTs), thousands of possible inputs, input routing takes many more bits than LUT contents:

16 bits to configure 4-input LUT, $\sim 16 \text{ bits} * 4$ to identify all inputs, in practice it takes ~ 50 bits per input to route each signal to the LUT.

- E.g., ~ 200 bits for inputs, 16 bits for logic

Real-Time Digital Simulation of the Human Nervous System (cont.) Engineering Risk Assessment

Uncertainty in synapse count:

Prefer Shepherd's numbers, \$ 85 M DRAM chip cost

Koch's numbers OK, \$170 M DRAM chip cost

Sejnowski's numbers considered high due to:

Perforated synapses: <http://synapses.mcg.edu/anatomy/radiatum/synapses.stm>

And multiple synapses per axon/dendrite pair

Biological unknowns:

Axon arbor sizes: 400 um axon & dendrite arbor diameter implies 50K axons within reach of each dendrite (assumes 100K incoming axons per mm² of cortex).

Adequate models of inputs and outputs (vision, hearing, musculature, etc)

Reliable operation:

600 Mhz DRAM data rate per pin

200 Mhz FPGA pipeline rate: faster chips coming (using 90nm fabs)

Reliable chips: neutron flux from cosmic rays may require shielding

<http://www.actel.com/products/rescenter/ser/docs/SERWP.pdf>

Routing and configuration:

Interconnect bandwidth & routing: 3D torus configuration has excess interconnect capacity

Placement: Distributing axon data across 20 parallel "processors" per FPGA

Use of Neural Models for Cognitive Processing (cont'd)

Digital/Hardware Learning

Hardware models of learning have characteristics not present in software models

Hardware is always active, e.g., computing

No penalty for parallelism

Must consider physical implementation of connectivity

Learning can take place via circuit design:

Valiant, LG (1984) *A theory of the learnable*, Communications of the ACM 27(11): 1134--1142.

Also consider decomposition theory:

B. Wurth, K. Eckl, and K. Antreich, *Functional multiple-output decomposition: Theory and an implicit algorithm*, in Proc. ACM/IEEE Design Automation Conf., June 1995, pp. 54--59.

Michael Burns, Marek Perkowski, and Lech Jozwiak, *An Efficient Approach to Decomposition of Multi-Output Boolean Functions with Large Set of Bound Variables*, in Proc. 1998 Euromicro, pp. 16-23, Vasteras, Sweden, August 25-27, 1998.

For practical work with learning via circuit design, take a VHDL or Verilog compiler (from Exemplar, Synplicity, Synopsys) and compile incompletely specified truth tables into ASIC gates, FPGA lookup tables or CPLD macro-cells.

Example of a real world implementation:

<http://www.dedicated-systems.com/vpr/layout/display/pr.asp?PRID=6802>

Real-Time Digital Simulation of the Human Nervous System (cont.)

Additional References

Many new books in Computational Neurobiology:

J. Feng ed, Computational Neuroscience: A Comprehensive Approach; 2002, CRC-Press.

Kotter ed, Neuroscience Databases: A Practical Guide; 2003, Kluwer Academic.

M.A. Arbib ed, The Handbook of Brain Theory and Neural Networks, 2nd ed; 2002, MIT Press.

P. Dayan & L.F. Abbott, Theoretical Neuroscience: Computational and Mathematical Modeling of Neural Systems; 2001, MIT Press.

F. Rieke, D. Warland, R.R. vanSteveninck & W. Bialek, Spikes: Exploring the Neural Code, 1999, MIT Press.

Ontogenesis:

D.H. Sanes, T.A. Reh & W.A. Harris, Development of the Nervous System, 2000, Academic Press.

Ultrastructure of the nervous system:

Josef Spacek: Atlas of Ultrastructural Neurocytology; <http://synapses.mcg.edu/index.asp>

Beginner's introduction to FPGAs & CPLDs, has history & glossary:


K. Parnell & N. Mehta: Programmable Logic Device Quick Start Handbook; 2002, PDF version available from Xilinx web site.

Large Scale Computing:

T. Sterling, P. Messina & P.H. Smith, Enabling Technologies for Petaflops Computing; 1995, MIT Press (the 10^6 FPGA system of this paper does $\sim 30 \text{ ops} * 40 \text{ Hertz} * 10^{14} \text{ synapses} = 120 \text{ Peta-ops}$).

Real-Time Digital Simulation of the Human Nervous System (cont.)

Golgi Atlas of the Postnatal Mouse Brain: www.cajal.csic.es/valverde



Real-Time Digital Simulation of the Human Nervous System (cont.)

Golgi Atlas of the Postnatal Mouse Brain: www.cajal.csic.es/valverde

